

# Analysis of the DNS Cache-Poisoning Attack using PRISM Model Checker



Academic Year 2025/2026

Giacomo Lombardi

Nicolò Zarulli

# Table of Contents

Introduction

Modelling the Kaminsky Attack on PRISM

Modelling DNS-CPM R<sub>ℓ1</sub> on PRISM

Modelling DNS-CPM R<sub>ℓ2</sub> on PRISM

Conclusion

# Table of Contents

## Introduction

Domain Name System (DNS)

DNS Cache-Poisoning Attacks

Mitigation:  
Randomization of  
Query ID and  
Source Port

Types of DNS  
Cache-Poisoning  
Attacks



# Table of Contents

## Modelling the Kaminsky Attack on PRISM

Transition Systems  
&  
Continuous Time Markov  
Chain

PRISM Model of  
Kaminsky DNS Cache-  
Poisoning Attack with  
Randomization of Query  
ID and Source Port

Property verification and  
experiments

DNSSEC solution, why  
not enough?



# Table of Contents

## Modelling DNS-CPM R<sub>l1</sub> on PRISM

DNS-CPM Mitigation

DNS-CPM Mitigation  
(R<sub>l1</sub>)

PRISM Model of  
Kaminsky DNS  
Cache-Poisoning  
Attack with DNS-CPM  
Mitigation

Property verification  
and experiments



# Table of Contents

## Modelling DNS-CPM RI2 on PRISM

PRISM Model of S<sub>Frag</sub>  
DNS Cache-Poisoning  
Attack

Property verification  
and experiments

DNS-CPM Mitigation  
(RI2)

PRISM Model of S<sub>Frag</sub>  
DNS Cache-Poisoning  
Attack with DNS-CPM  
Mitigation

Property verification



# Table of Contents

## Conclusion

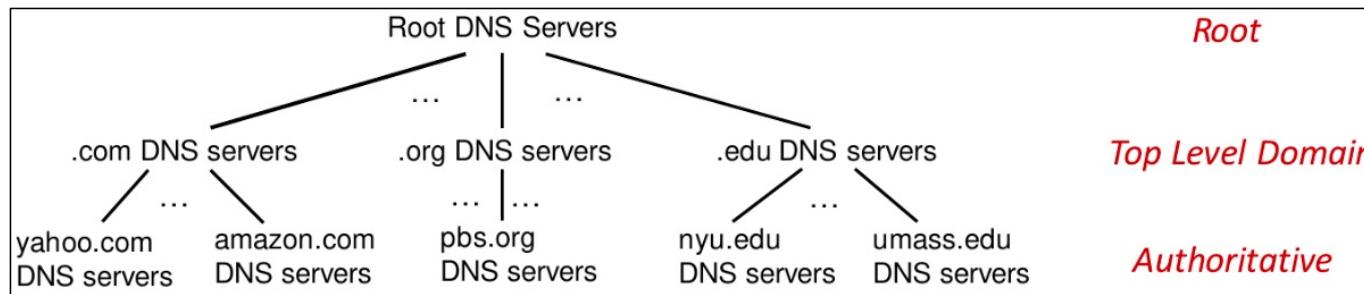
Summarizing

Observations  
and  
conclusions



# Domain Name System (DNS)

- DNS is a distributed database with hierarchical structure

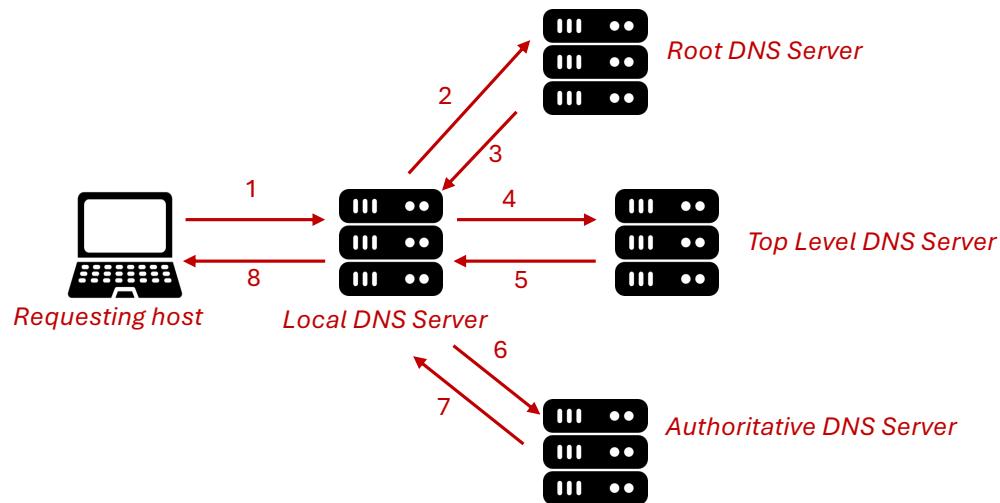


- An entry in the database is called **Resource Record**:  $\langle \text{Name}, \text{Value}, \text{Type}, \text{Class}, \text{TTL} \rangle$ 
  - *Name* and *Value* depend on the *Type*
  - Most famous types are: “A”, corresponding to the mapping of name-IPv4, and “AAAA”, corresponding to the mapping name-IPv6
  - *TTL* specifies how long the RR is valid (can be kept in cache)
- DNS is an application layer protocol
- Most DNS traffic is over **UDP** since DNS is just a request/replay protocol



# Domain Name System (DNS)

- DNS name resolution approach: **Iterative**



## Example

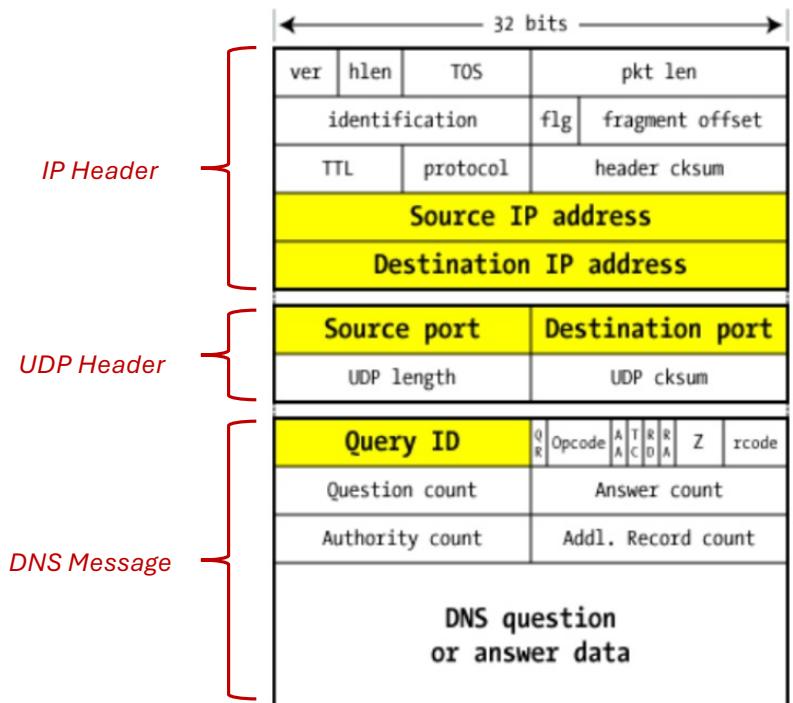
1. Host → Resolver: IP of www.google.com
2. Resolver → Root DNS Server: IP of www.google.com
3. Root DNS Server → Resolver: IP of TLD DNS Server (.com DNS Server)
4. Resolver → TLD DNS Server: IP of www.google.com
5. TLD DNS Server → Resolver: IP of Authoritative DNS Server (google.com DNS Server)
6. Resolver → google.com DNS Server: IP of www.google.com
7. Google.com DNS Server → Resolver: Authoritative Answer (AA) containing the IP
8. Resolver → Host: forwards AA

- Local DNS Server = Resolver → extracts information from DNS Servers in response to clients' requests
- Resolver saves Domain Name-IP associations in its **cache**



# Domain Name System (DNS)

- DNS packet
- DNS message encapsulated in an UDP header
- Which is encapsulated in an IP header
- DNS query:
  - **Source IP** = IP of the client sending the DNS query
  - **Destination IP** = IP of the server
  - **Source port** = chosen by the OS
  - **Destination port** = 53
  - Addresses are exchanged in the response
  - **Query ID**: unique identifier created in the query packet; the responding server uses the same Query ID



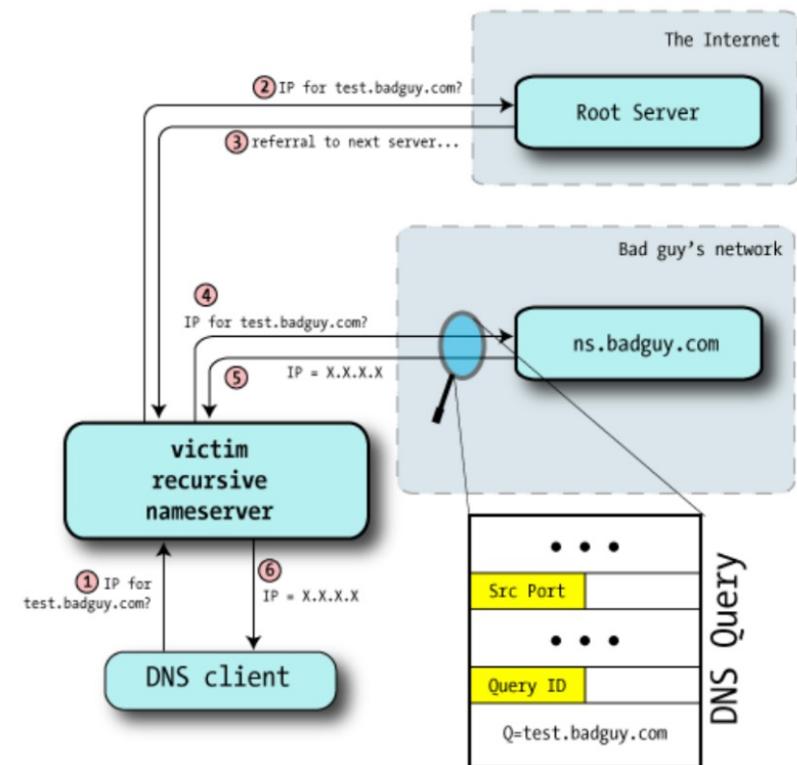
# DNS Cache-Poisoning Attacks

**Idea:** the attacker inserts wrong information in the cache of the resolver, so that if some user makes a certain DNS query will receive a wrong response.

**Intuitively:** When the resolver sends a request to an Authoritative DNS Server, if the attacker can forge a formally correct DNS response, the resolver will accept it as correct if it arrives before the legitimate one from the Authoritative DNS Server.

**Formally correct DNS response:** The formally correct DNS response has the same Query ID of the DNS request and Destination Port equal to the Source Port of the DNS request

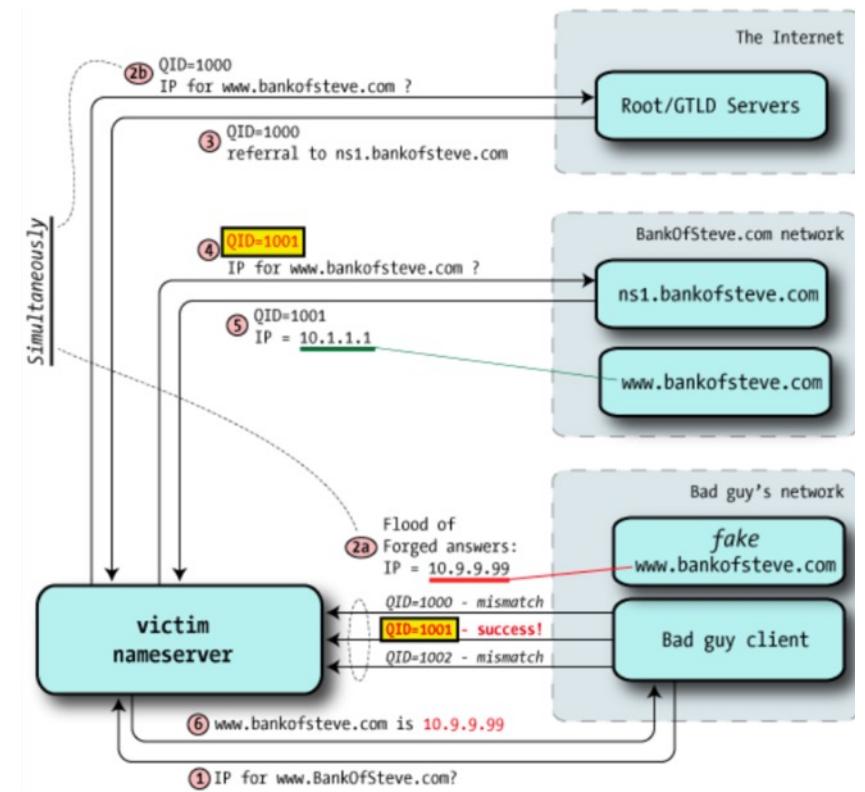
**Original vulnerabilities:** {  
    Resolver always use the same Source Port  
    Incremental Query ID for each new request made by the resolver



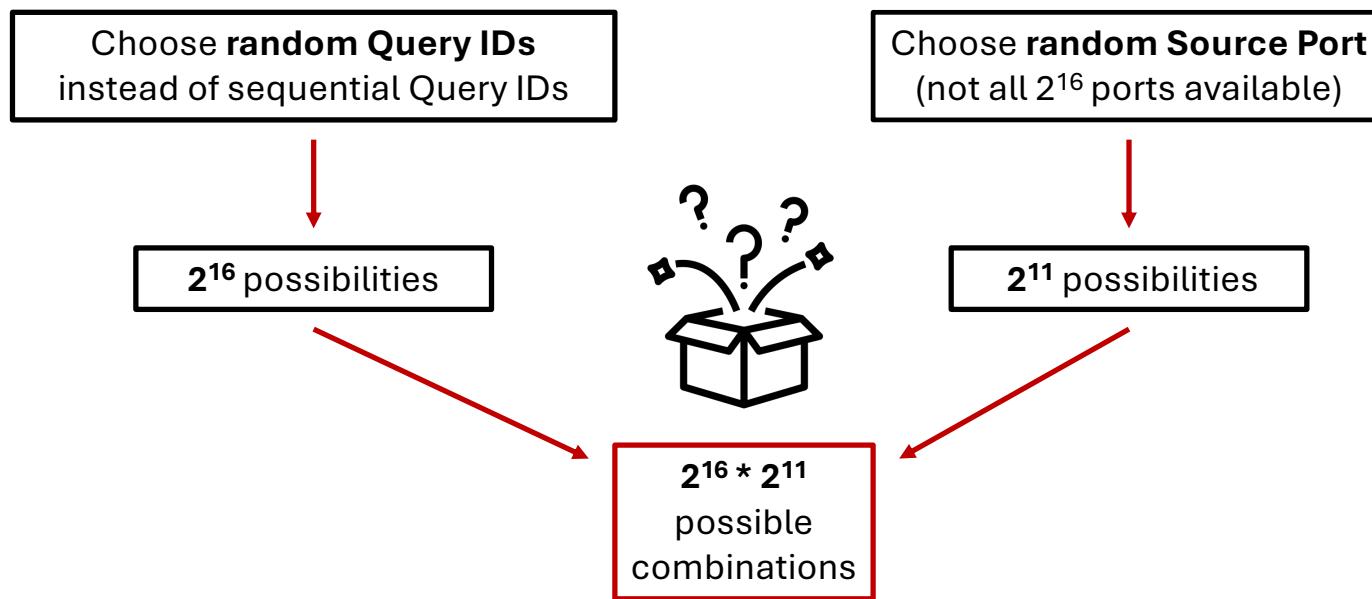
# DNS Cache-Poisoning Attacks

## General attack steps:

1. Suppose the attacker has obtained the Source Port used by the resolver to make queries and a reference value for Query ID.
2. Trick the resolver to query the Authoritative DNS Server of the target domain  $D$ , expecting a response with the IP address of  $D$ .
3. While the resolver is expecting a response from authoritative DNS Server, attacker provides it with a forged response with the source IP of the Authoritative DNS Server.
4. For the forged response to be accepted by the resolver it must have the correct Destination Port number, Query ID, and domain  $D$ , in addition to the correct IP address of Authoritative DNS Server.



## Mitigation: Randomization of Query ID and Source Port



# Types of DNS Cache-Poisoning Attacks



**S type (Kaminsky attack):** Attacker floods the resolver with many guesses trying different combinations of Destination Port number and Query ID, hoping that one of them will have the correct numbers. It has more or less  $2^{32}$  possible combinations.



**S<sub>Frag</sub> type (Fragmentation attack):** When fragmentation of IP packets is used, second fragments lack UDP and DNS headers, meaning no Destination Port and no Query ID have to be guessed. The only security check is the 16-bit IPID ( $2^{16}$  possibilities).

Attacker first sends a series of forged *2<sup>nd</sup> fragment*, with different IPIDs, containing the mapping it wants to be cached. Then, it tricks the resolver to query Authoritative DNS Server with a query whose response size requires fragmentation.

The *2<sup>nd</sup>* fragments wait in the OS reassembly buffer and when the legitimate *1<sup>st</sup>* fragment arrives, the OS reassembles it with the malicious *2<sup>nd</sup>* fragment if the IPID matches



**S<sub>OoB</sub> type (Out-of-Bailwick):** This is a variant of Kaminsky attack, in which the attacker includes in the forged response an OoB record with the mapping it wishes to insert into the cache (e.g. a mapping for *example.net* while the original query was for *bankofsteve.com*. The *.net* domain is out of bailwick for *.com*). Due to inadequate validation, the resolver may cache the additional OoB mappings.



# Transition Systems

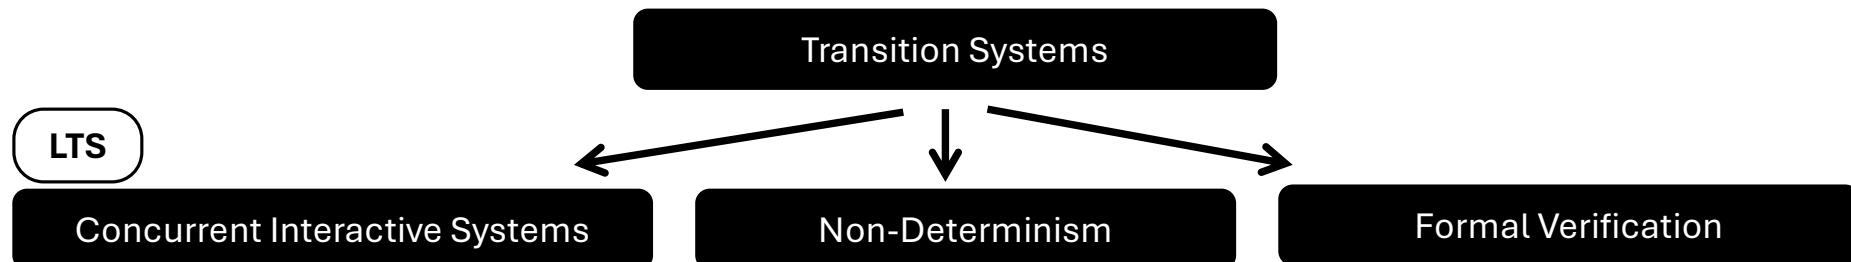
The DNS Cache Poisoning Attacks works on the basis of a



**Time Based Race Condition**

Intruder vs. Authoritative Name Server

We want to model the environment in which an attacker can pose a DNS Cache Threat and being able to trace all the possible paths that may lead to an attack. Given therefore that we are interested in studying whether a Victim DNS Resolver reaches a bad state, hence “cache poisoned”, **Transition Systems** offer us the underlying structure to do just that.



We need to model an environment made of different **independent components** that may **perform actions synchronizing** with each other and the environment

To model choices and uncertainty

We use PRISM both to model the system and to carry out a process of Model Checking



# Continuous Time Markov Chain

The DNS Cache Poisoning Attacks work on the basis of a



**Time Based Race Condition**

Intruder vs. Authoritative Name Server

The suitable choice for us to model the behavior of a DNS Cache Poisoning Attack is the **Continuous Time Markov Chain (CTMC)**. Why a CTMC instead of a DTMC?

## Real Time vs Step-Based Time



The DTMC progresses in fixed steps, and at each step one transition occurs. Almost like the ticking of a clock.



**Problem:** DNS packets won't wait for a clock tick but may be sent at any given time. In a CTMC a transition can happen at any moment (random intervals).

## Attacker Race Problem



The DTMC struggles with races, hence evaluating the next transition when two things are happening at the same time but at two different speeds.



The CTMC uses the concept of **Rates** to evaluate the "next transition". In the case that two events are scheduled, the one with the higher rate is statistically more likely to happen!

# Continuous Time Markov Chain

By using a CTMC we are then able to model:

## The Race Condition



Cache can be either «poisoned» or «safe»



### Winner Takes All

**Philosophy:** The first packet to arrive and be accepted determines the state (poisoned vs. safe).



In a CTMC the fastest transition determines the next state



## Real World Metrics and Attack Parameters



DNS Attack Traffic and Real Network Traffic are defined by **Rates (events per second)**.



CTMCs use a Transition Rate Matrix where rates are based on **real-world parameters** (Attacker Guess Rate vs DNS Server Response Arrival Rate).



The model can account for Server Load where higher traffic **slows** down the legitimate server, **increasing** the intruder's window of opportunity

## Queueing and Delays

### Waiting Times



Time spent in a state is governed by a negative exponential distribution, simulating variable network latency and processing delays.

### Vulnerability Window



The model captures the specific duration the victim server is "waiting" for a reply, which is the exact window the attacker must exploit

# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with Randomization of Query ID and Source Port

## Modules

 **Intruder Machine**  
Intruder's PC that requests resolution of a URL

 **Victim Resolver**  
Local DNS Server that responds to the client's requests

**Root DNS Server**  
DNS Server that knows the IP address of the Authoritative DNS Server

**Domain Server**  
Authoritative DNS Server of the requested domain

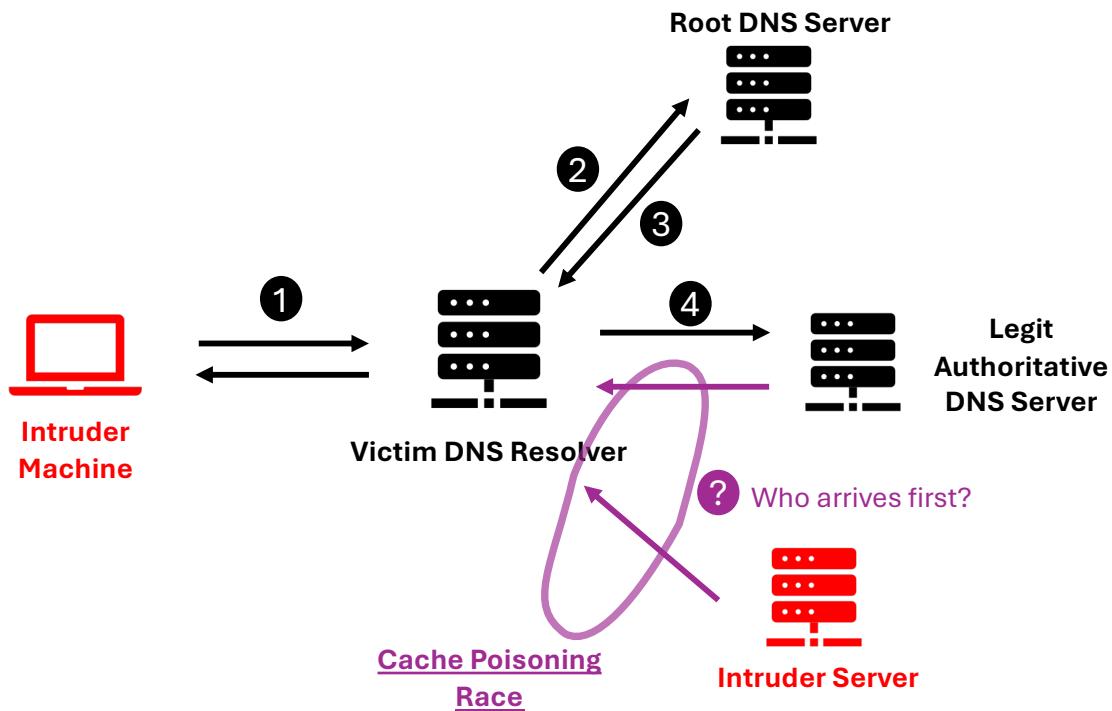
 **Intruder Server**  
Intruder's DNS Server trying to guess correct Query ID (and Port number)

## Final states of the model

- 1) A correct guess has been made → Cache-Poisoning Attack success
- 2) All requests from Intruder Machine answered correctly



## PRISM Model of Kaminsky DNS Cache-Poisoning Attack with Randomization of Query ID and Source Port



# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with Randomization of Query ID and Source Port

## Model parameters

-  **Number of URL requests:** Number of URL requests attacker sends to resolver (i.e. different attacks)
-  **Max queue:** Total number of requests attacker can send to resolver
-  **Popularity:** Rate at which the TTL of the resolver's cache entry for the domain the attacker is trying to poison has a positive value
-  **Query ID range** Number of possible Query IDs (set by default to  $2^{16} = 65536$ )
-  **Port ID range:** Range of Port numbers for the purpose of implementing source-port randomization
-  **Guess:** Rate at which the attacker server sends fake responses to resolver
-  **Requests rate:** Rate at which the attacker sends requests to resolver
-  **Authoritative DNS Workload:** Workload on the Authoritative DNS Server for the “attacked” domain, related to traffic different than the one from victim resolver



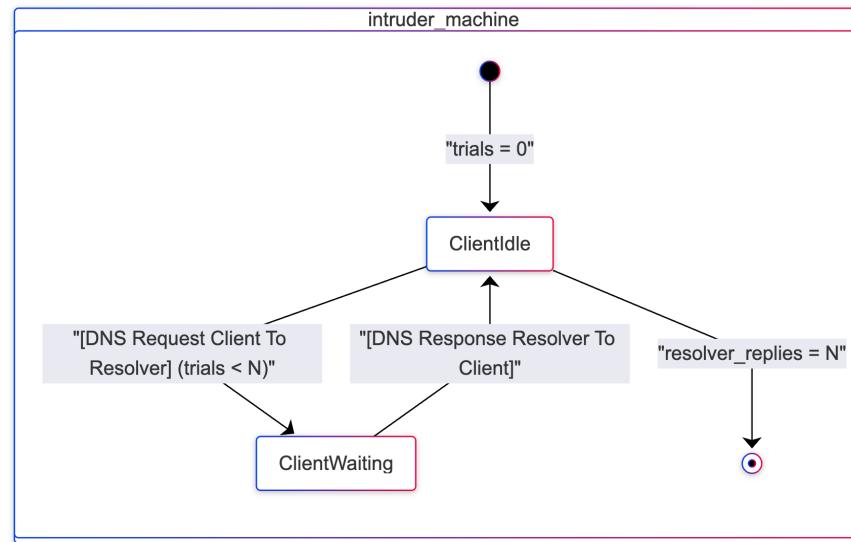
# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with Randomization of Query ID and Source Port

```
module intruder_machine
    trials : [0..NUMBER_OF_URL_REQUESTS] init 0;
    resolver_replies : [0..NUMBER_OF_URL_REQUESTS] init 0;

    // -- Transitions
    // 1. Send Request: Initiates resolution.
    [DNS_Request_Client_To_Resolver] trials < NUMBER_OF_URL_REQUESTS -> requests_rate : (trials' = trials + 1);

    // 2. Receive Reply: Marks end of query cycle.
    [DNS_Response_Resolver_To_Client] resolver_replies < trials -> (resolver_replies' = resolver_replies + 1);

    // 3. Termination
    [] resolver_replies = NUMBER_OF_URL_REQUESTS -> true;
endmodule
```



# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with Randomization of Query ID and Source Port

```
module victim_resolver
  ttl : [0..2] init 2;
  client_queue : [0.. MAX_QUEUE] init 0;
  root_queue : [0.. MAX_QUEUE] init 0;
  domain_queue : [0.. MAX_QUEUE] init 0;
  responses_queue : [0.. MAX_QUEUE] init 0;
  query_to_root_server: bool init false;
  query_to_domain_server: bool init false;
  answer_from_domain_received: bool init false;
  correct_guess: bool init false;

  // -- Transitions
  // 1. Client Request (Merged Hit/Miss Logic)
  [DNS_Request_Client_To_Resolver] (client_queue < MAX_QUEUE) & (responses_queue < MAX_QUEUE) & (root_queue < MAX_QUEUE) ->
    // Branch 1: Cache Hit
    (popularity / 10) : (client_queue' = client_queue + 1) & (ttl' = 1) & (responses_queue' = responses_queue + 1)
    +
    // Branch 2: Cache Miss
    (1 - (popularity / 10)) : (client_queue' = client_queue + 1) & (ttl' = 0) & (root_queue' = root_queue + 1);

  // 2. Query Root
  [DNS_Request_Resolver_To_Root] (client_queue > 0) & (root_queue > 0)
    -> (query_to_root_server' = true);

  // 3. Root Responds -> Move to Domain Queue
  [DNS_Response_Root_To_Resolver] (root_queue > 0) & (domain_queue < MAX_QUEUE)
    -> (root_queue' = root_queue - 1) & (domain_queue' = domain_queue + 1) & (query_to_root_server' = false);

  // 4. Query Domain (Opens Vulnerability Window)
  [DNS_Request_Resolver_To_Domain] domain_queue > 0 -> (query_to_domain_server' = true);

  // 5. Domain Responds (Closes Window)
  [DNS_Response_Domain_To_Resolver] (domain_queue > 0) & (responses_queue < MAX_QUEUE) & (correct_guess = false)
    -> (domain_queue' = domain_queue - 1) & (responses_queue' = responses_queue + 1) & (query_to_domain_server' = false)
    & (answer_from_domain_received' = true);]

  // 6. Reply to Client
  [DNS_Response_Resolver_To_Client] (responses_queue > 0) & (client_queue > 0)
    -> (client_queue' = client_queue - 1) & (responses_queue' = responses_queue - 1);
```



# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with Randomization of Query ID and Source Port

```
// --- Race Condition (Attack) ---
[Guess] (correct_guess = false) & (query_to_domain_server = true) ->
    1 / (query_id_range * port_id_range) : (correct_guess'=true)
    +
    ((query_id_range * port_id_range) - 1) / (query_id_range * port_id_range): (correct_guess'=false);

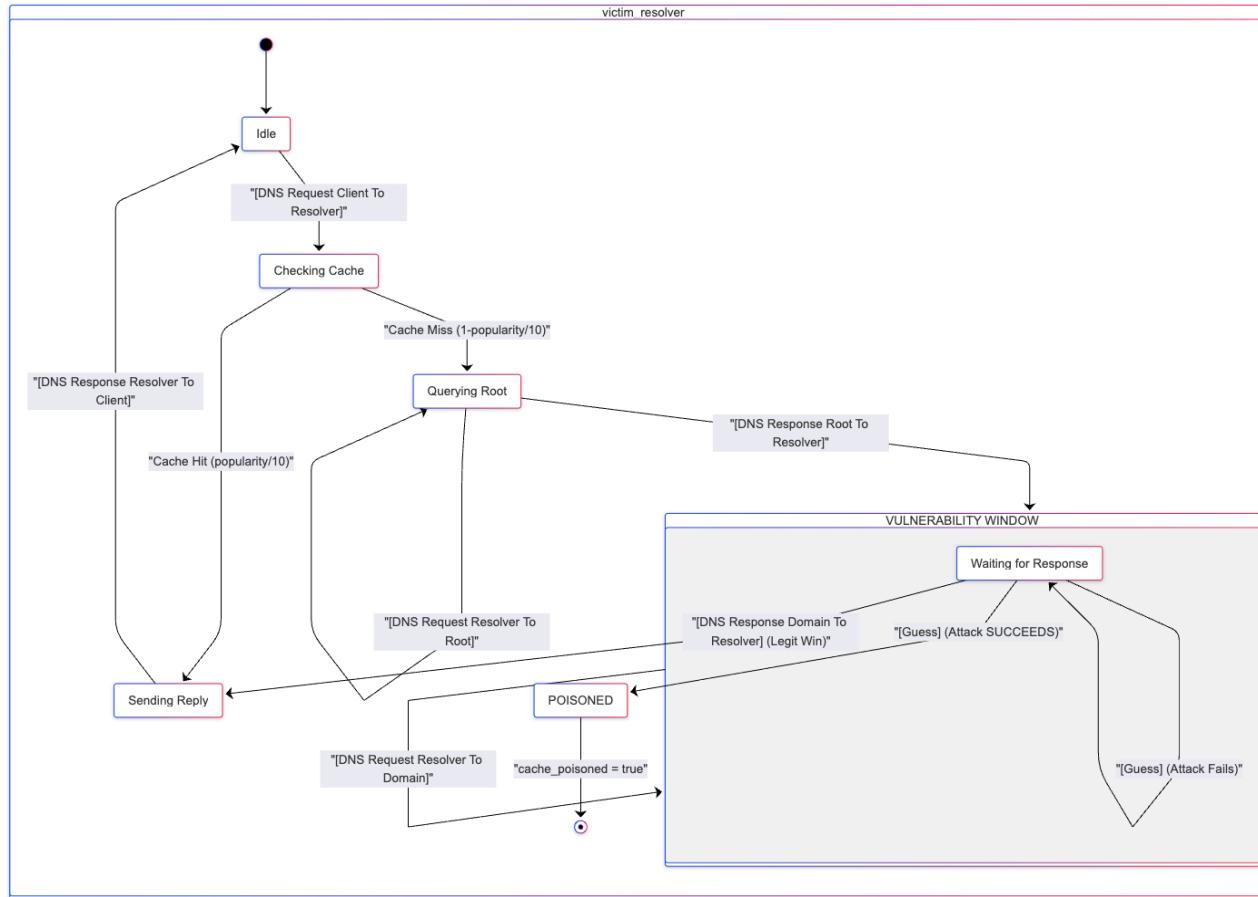
// Sink State
[] cache_poisoned = true -> true;
endmodule
```

Rate with which the attacker makes a correct guess:

$$guess * \frac{1}{query\_id\_range * port\_id\_range} = guess * \frac{1}{2^{16} * port\_id\_range}$$



# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with Randomization of Query ID and Source Port



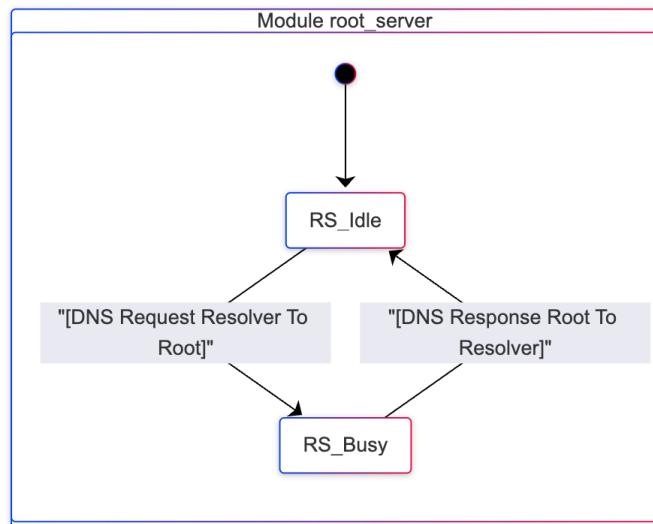
# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with Randomization of Query ID and Source Port

```
module root_server
    root_state : [0..1] init 0;

    // Transitions

    // 1. Receive Request: A query arrives from the Victim Resolver.
    [DNS_Request_Resolver_To_Root] root_state = 0 -> (root_state' = 1);

    // 2. Send Reply: The Root Server immediately responds, moving the resolver to the next step (Domain Server).
    [DNS_Response_Root_To_Resolver] root_state = 1 -> (root_state' = 0);
endmodule
```

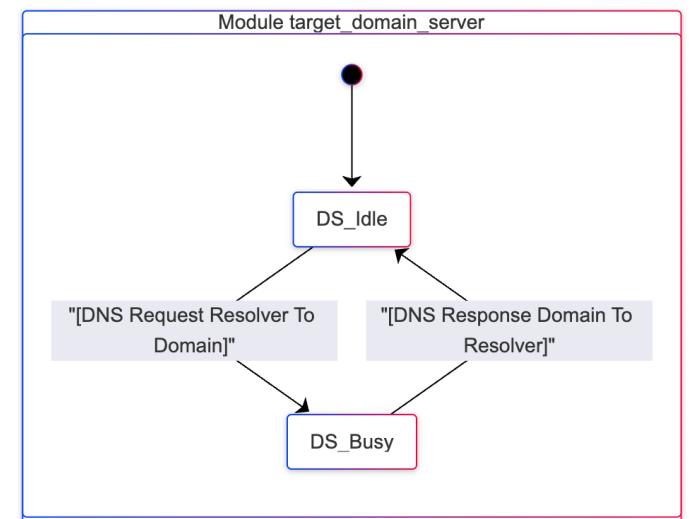


# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with Randomization of Query ID and Source Port

```
module target_domain_server
  // -- States
  domain_state : [0..1] init 0;

  // -- Transitions
  // 1. Receive Request: The Victim Resolver asks "Where is example.com?".
  // This opens the 'Race Window'. The Intruder starts guessing NOW.
  [DNS_Request_Resolver_To_Domain] domain_state = 0 -> (domain_state' = 1);

  // 2. Legitimate Response (The "Good" Ending):
  // The server processes the specific request from our Victim Resolver.
  // This transition closes the 'Race Window' (Intruder stops guessing).
  [DNS_Response_Domain_To_Resolver] domain_state = 1 -> 1/(authoritative_dns_workload) : (domain_state' = 0);
endmodule
```



# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with Randomization of Query ID and Source Port

```

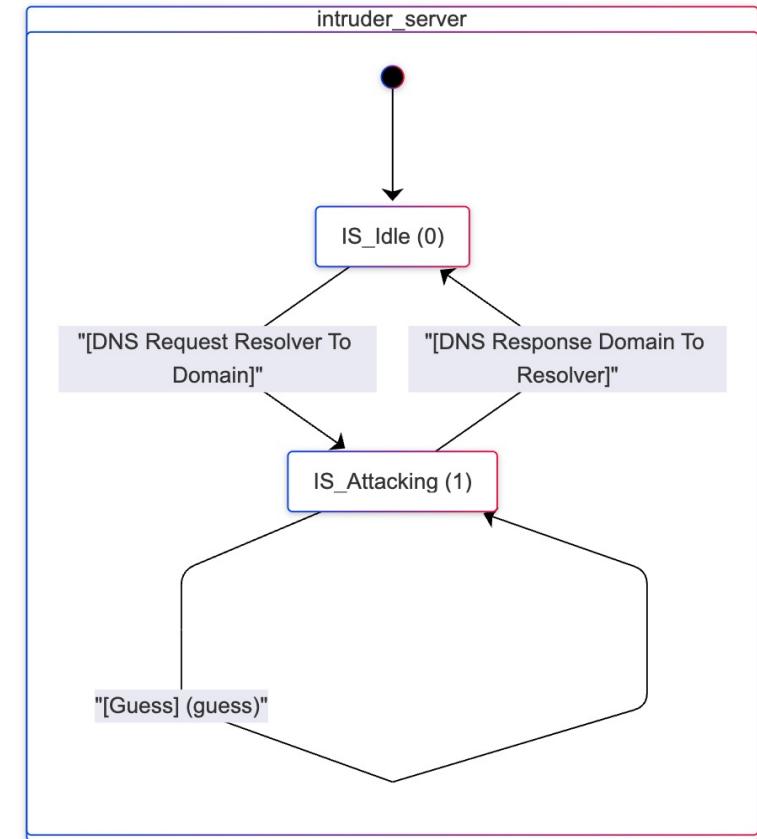
module intruder_server
    is_state : [0..1] init 0;

    // -- Transitions
    // 1. Start Attack: Triggered when the Victim Resolver queries the Domain Server.
    // The 'Race Window' is now open.
    [DNS_Request_Resolver_To_Domain] is_state = 0 -> (is_state' = 1);

    // 2. Firing Loop (The "Machine Gun"):
    // Fires guesses at rate 'guess'.
    // The success/failure logic is handled by the Victim Resolver module via probabilistic split.
    [Guess] is_state = 1 -> guess : (is_state' = 1);

    // 3. Stop Attack: Triggered when the Victim Resolver receives the legitimate answer.
    // The 'Race Window' is closed. The attack failed for this specific query.
    [DNS_Response_Domain_To_Resolver] is_state = 1 -> (is_state' = 0);
endmodule

```



# Property verification and experiments

P = ? [F cache\_poisoned]

## Property Verification

**Property:**

P=? [ F cache\_poisoned ]

**Defined constants:**

NUMBER\_OF\_URL\_REQUESTS=1,popularity=0,port\_id\_range=1,guess=300,authoritative\_dns\_workload=150

**Method:**

Verification

**Result (probability):**

0.4071071867990663 (value in the initial state)



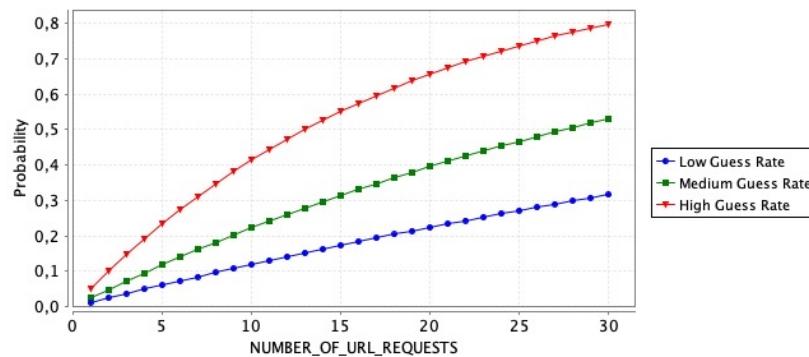
With only Query ID randomization (no Port randomization), a high guess rate, and a high workload on the Authoritative DNS Server, the attacker has a **40%** of ***probability of poisoning the cache*** of the resolver



# Property verification and experiments

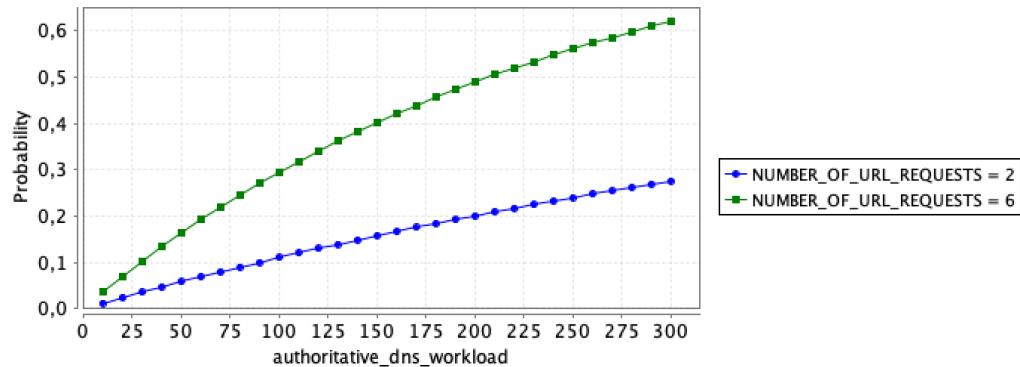
P = ? [F cache\_poisoned]

## Experiments

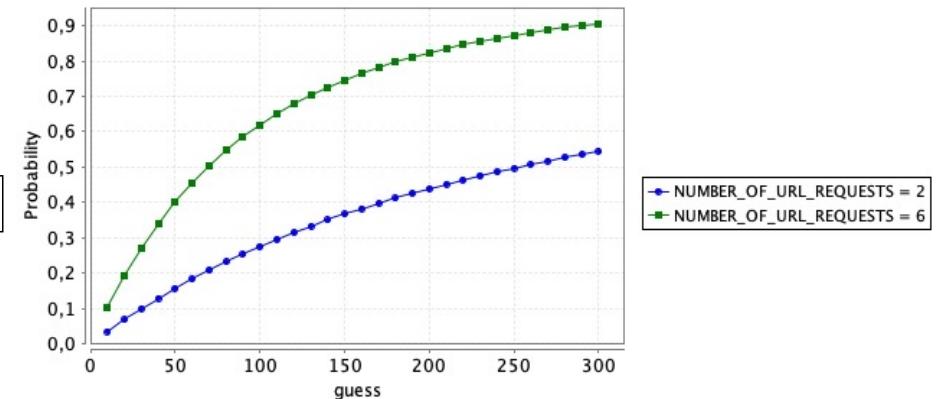


For three values of parameter guess (low=30, medium=60, high=130), we vary *number\_of\_url\_requests* from 1 to 30; *authoritative\_dns\_workload*=50 and *port\_id\_range*=1. As expected, the more URL-resolution requests there are for the target domain, the higher the probability of attack success

# Property verification and experiments



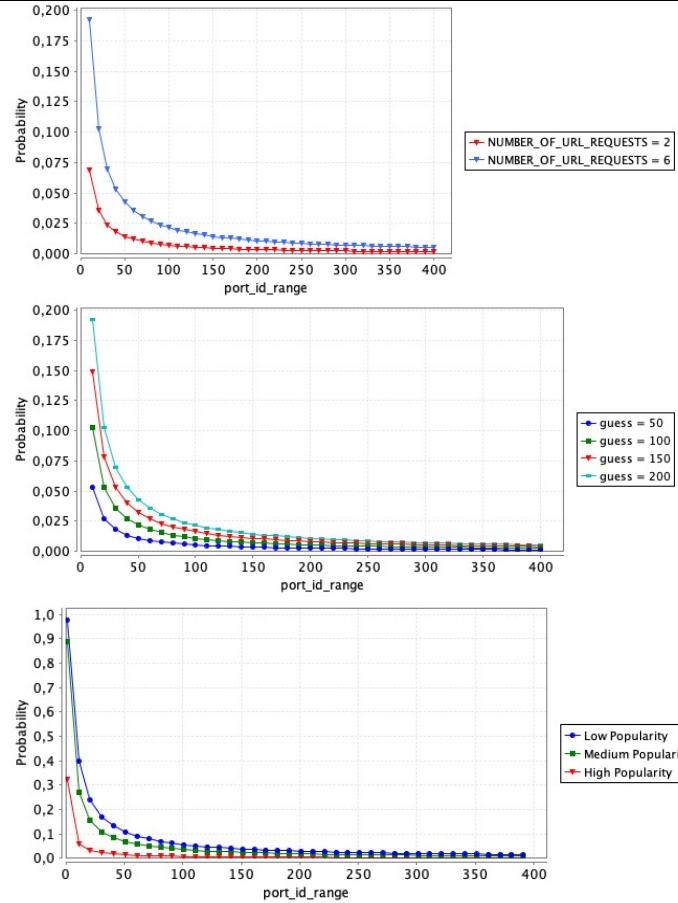
For two values of parameter *number\_of\_url\_requests* (2 and 6), we vary *authoritative\_dns\_workload* from 1 to 300; *guess*=50 and *port\_id\_range*=1.  
As expected, the higher the workload on Authoritative DNS Server, the higher the probability of attack success.



For two values of parameter *number\_of\_url\_requests* (2 and 6), we vary *guess* from 1 to 300; *authoritative\_dns\_workload* =150 and *port\_id\_range*=1.  
As expected, increasing rate *guess* increases the probability of attack success



# Property verification and experiments

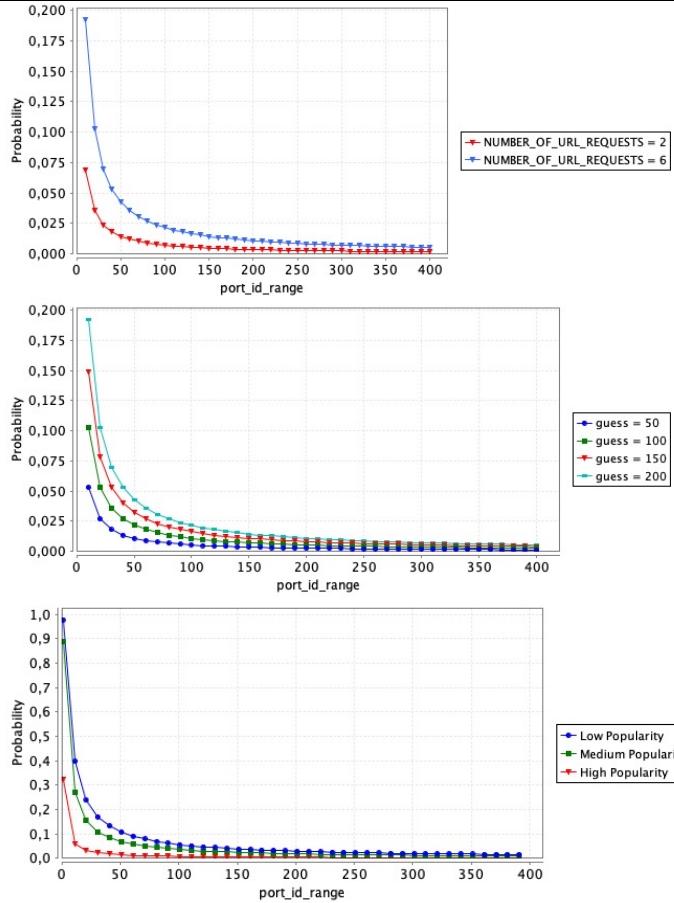


For two values of parameter `number_of_url_requests` (2 and 6), we vary `port_id_range` from 1 to 400; `authoritative_dns_workload` =150 and `guess`=200.

Considering four values of parameter `guess` (50, 100, 150, and 200), we vary `port_id_range` from 1 to 400; `number_of_url_requests`=6 and `authoritative_dns_workload` =150.

For three values of parameter `popularity` (`low`=2, `medium`=5, `high`=9), we vary `port_id_range` from 1 to 400; `authoritative_dns_workload` =300 and `number_of_url_requests`=8.

## Property verification and experiments

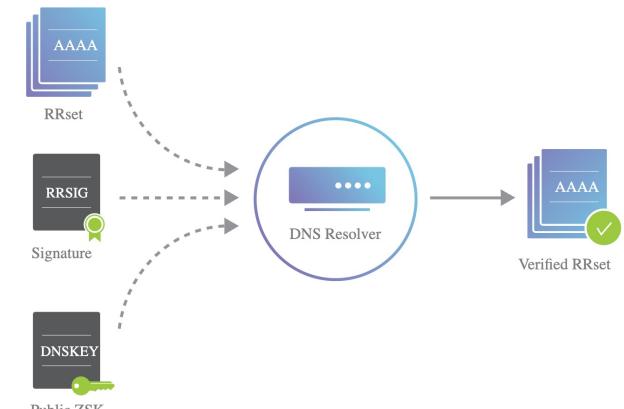
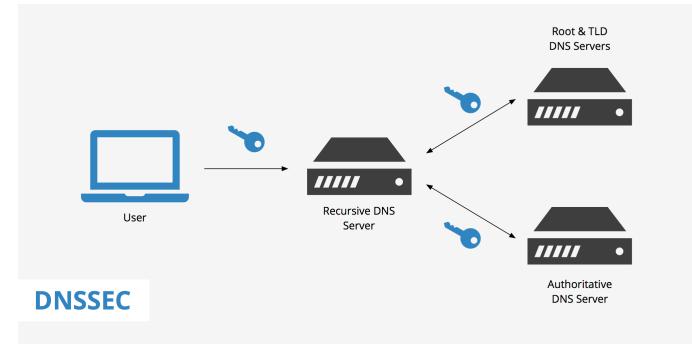


The probability of attack success decreases nonlinearly as the value of *port\_id\_range* increases, because the attacker has to guess both Query ID and Port Number



# One Possible Solution: DNSSEC

- 💡 **DNS Security Extension (DNSSEC)** protocol was designed to eliminate DNS Cache-Poisoning Attacks exploiting *digital signatures*.
- 🛡️ Guarantees *Integrity* and *Authentication* of DNS RR
- 🛡️ Protects from Spoofing, MIM Attacks and Cache Poisoning
- 🤝 Establishes a «**Chain of Trust**» based on Zone Authority that authenticates DNS Resource Records from their origin, given that DNS resolves queries with an *iterative* approach.
- 📦 **New DNS Resource Records:** RRSIG, DNSKEY, DS, NSEC, CDNSKEY
- ✍️ **Signature of a RR Set:** Multiple RR Authenticated at the same time. Usage of Asymmetric Zone Keys (ZSK).



# DNSSEC solution: why not enough?

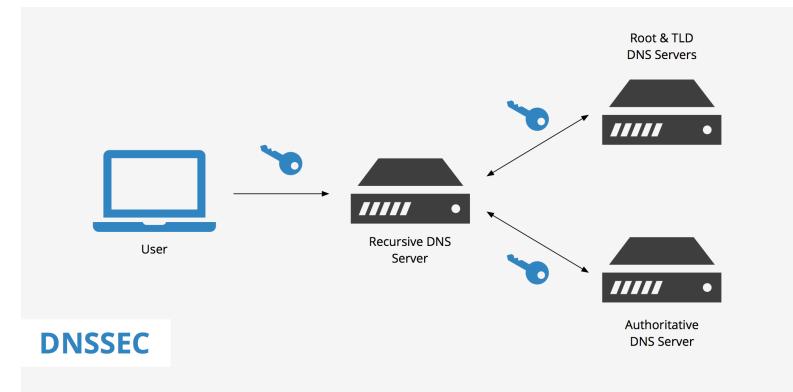
**DNS Security Extension (DNSSEC)** protocol was designed to eliminate DNS Cache-Poisoning Attacks exploiting ***digital signatures***.

However, we cannot only rely on DNSSEC for the following reasons:

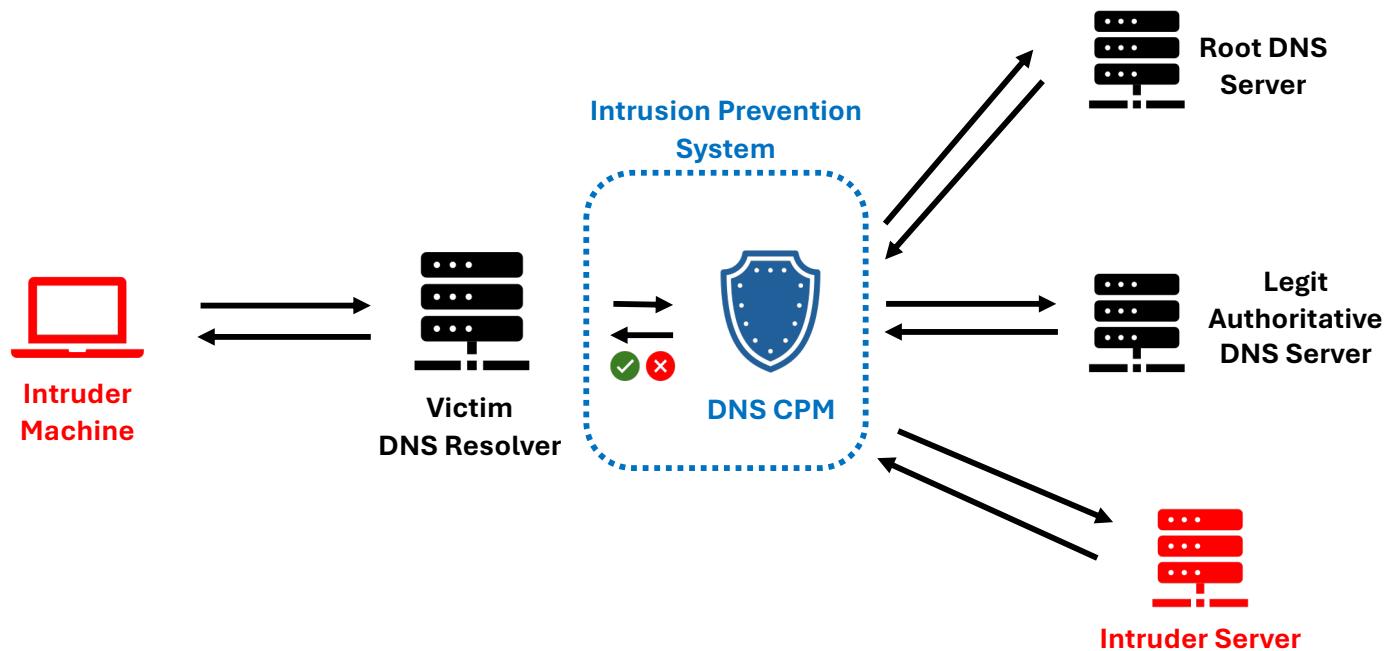
- 1) Its adoption is ***limited*** (<30% of DNS Servers over the Internet)
- 2) It heavily ***depends on*** Authoritative DNS Servers implementation



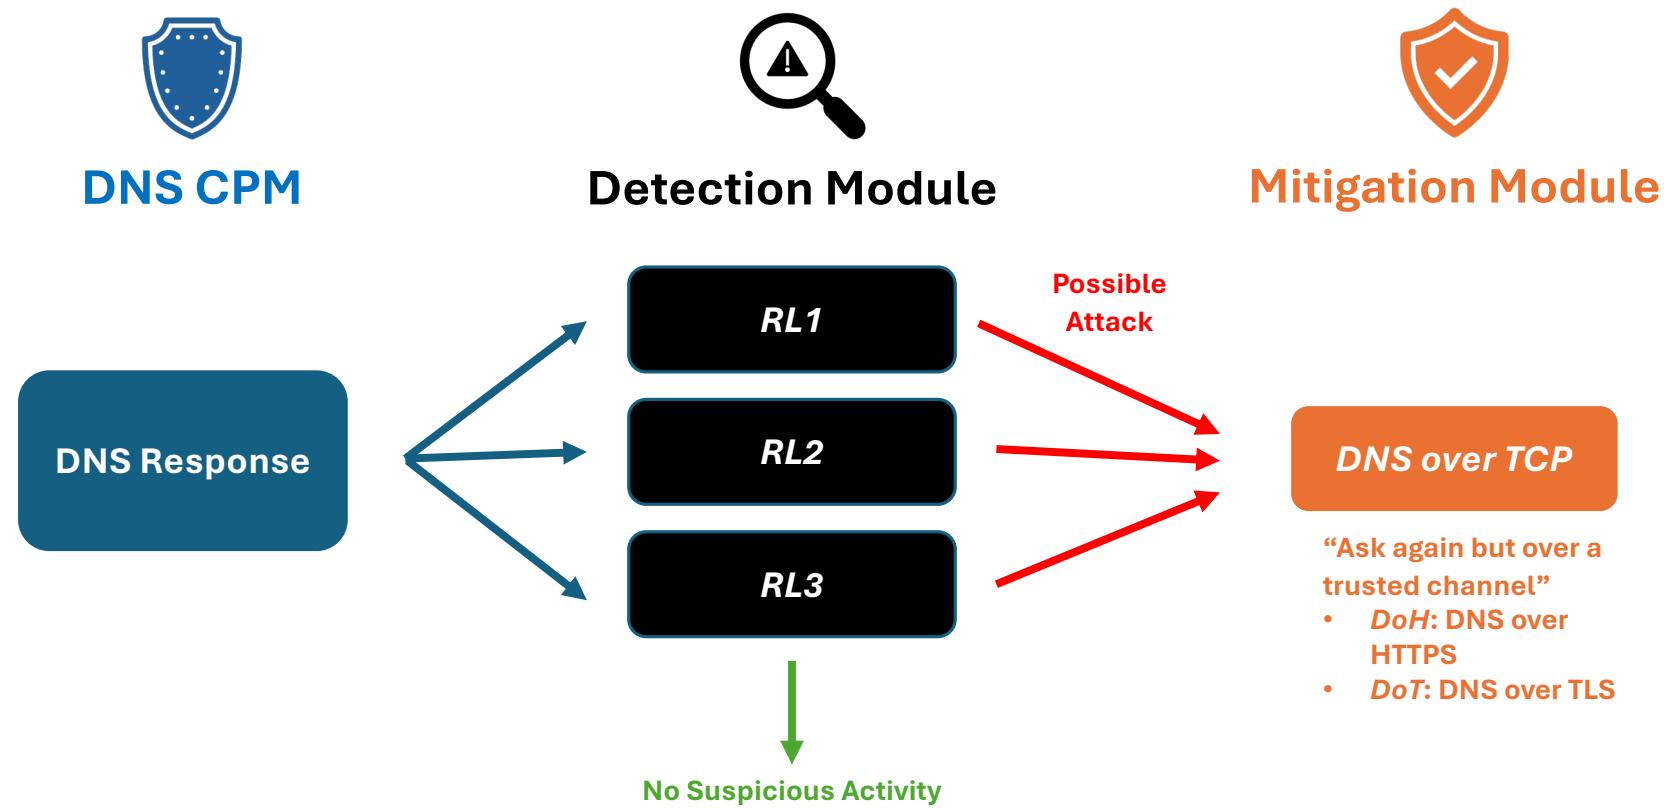
A possible **solution** that moves protection to the resolver is **DNS-CPM**, which acts as an *Intrusion Prevention System (IPS)* against DNS Cache-Poisoning Attacks



# DNS-CPM



# DNS-CPM Detection & Mitigation



# DNS-CPM Detection & Mitigation



## Detection Module



Applies different rules to identify different types of DNS Cache-Poisoning Attacks



**R1** against  $S$  attacks, **R2** against  $S_{Frag}$  attacks, and **R3** against  $S_{OoB}$  attacks



When responses are considered suspicious, they are flagged as **potential poisoning attack** and passed to the Mitigation Module



## Mitigation Module



When potential poisoning attack responses arrive, the TC flag of those packets is set to true.



When a packet with **TC = true** arrives to the resolver, query/response conversation is moved from DNS over UDP to **DNS over TCP**.



Even if a correct guess arrives to the resolver (matching both Query ID and Port number), it opens a TCP session with authoritative DNS Server, thus preventing the attack



# DNS-CPM Mitigation ( $R\ell 1$ )



## Detection Module ( $R\ell 1$ )



Monitor the number of DNS response packets for a URL request that differ from each other only in one common field, i.e. destination port or Query ID.

To do so, we considered the **Count-Min Sketch (CMS)** algorithm (error rate = 0.01%)



When the counted number crosses a *threshold* (e.g. 5 packets), the responses are flagged as a **potential poisoning attack** and passed to the Mitigation Module



# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with DNS-CPM Mitigation

## Modules

 **Intruder Machine**  
Intruder's PC that requests resolution of a URL

 **Victim Resolver**  
Local DNS Server that responds to the client's requests

**Root DNS Server**  
DNS Server that knows the IP address of the Authoritative DNS Server

**Domain Server**  
Authoritative DNS Server of the requested domain

 **Intruder Server**  
Intruder's DNS Server trying to guess correct Query ID (and Port number)



 **DNS-CPM Module**  
IPS that implements Detection with R&I and Mitigation modules



# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with DNS-CPM Mitigation

## Additional model parameters

5

**Tau**

Threshold of DNS response packets for a URL request with identical fields, except for Query ID/Port number. When the counter in the DNS-CPM Module exceeds Tau, the mitigation mechanism is activated



**CMS error rate**

Considering the way CMS works, this represents the rate of other responses that arrive to the victim resolver that are wrongly associated to the request made by the attacker



CMS uses *hash functions* to verify if response packets differ only for one field value. Collisions' probability is low, but not zero.



**Benign noise rate**

Network traffic arriving to the resolver other than the one related to the attacker.

It represents legitimate responses to requests of other clients.



# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with DNS-CPM Mitigation

```

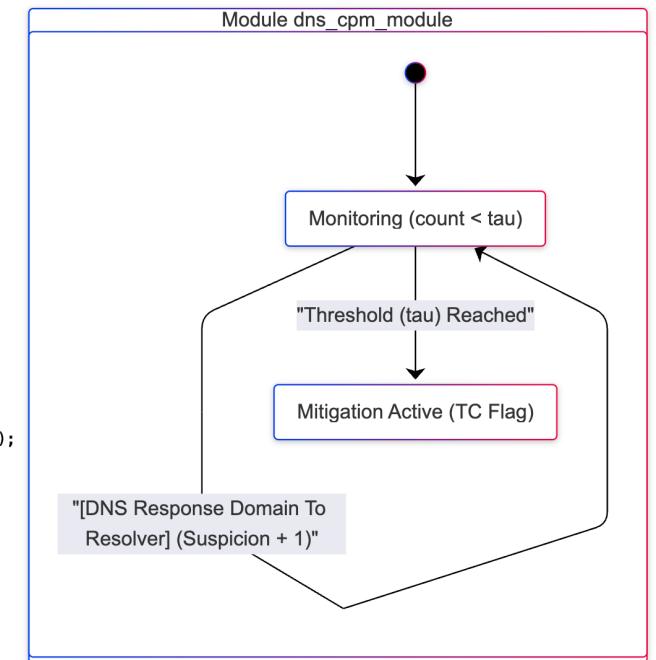
module dns_cpm_module
    suspicious_count : [0..tau] init 0;      // Count-Min Sketch counter
    detection_signal : bool init false;        // Signal to the Mitigation Module (True if threshold reached)

    // 1. Detection of a Possible Attack
    [Guess] suspicious_count < tau ->
        (suspicious_count' = suspicious_count + 1) &
        (detection_signal' = (suspicious_count + 1 >= tau));

    // 2. Detection of Benign Noise: Normal traffic towards the victim resolver. If a hash collision occurs, increment counter.
    [Benign_Response] suspicious_count < tau ->
        // Case 1: Collision occurs
        cms_error_rate * ((benign_noise_rate - 1) / benign_noise_rate) :
            (suspicious_count' = suspicious_count + 1) & (detection_signal' = (suspicious_count + 1 >= tau))
        +
        // Case 2: No Collision
        (1 - cms_error_rate) * (1 / benign_noise_rate) : (suspicious_count' = suspicious_count) & (detection_signal' = detection_signal);

    // 3. Detection of a Legitimate Domain Response: Counts towards the threshold (Same Domain)
    [DNS_Response_Domain_To_Resolver] suspicious_count < tau ->
        (suspicious_count' = suspicious_count + 1) &
        (detection_signal' = (suspicious_count + 1 >= tau));
endmodule

```



# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with DNS-CPM Mitigation

```

module victim_resolver
    ttl : [0..2] init 2; // 0=Miss, 1=Hit, 2=Init
    client_queue : [0..MAX_QUEUE] init 0;
    root_queue : [0..MAX_QUEUE] init 0;
    domain_queue : [0..MAX_QUEUE] init 0;
    responses_queue : [0..MAX_QUEUE] init 0;

    // -- Variables
    query_to_root_server: bool init false;
    query_to_domain_server: bool init false;
    answer_from_domain_received: bool init false;
    correct_guess: bool init false;

    // -- Transitions
    // 1. Client Request (Merged Hit/Miss Logic)
    [DNS_Request_Client_To_Resolver] (client_queue < MAX_QUEUE) & (responses_queue < MAX_QUEUE) & (root_queue < MAX_QUEUE) ->
        // Branch 1: Cache Hit
        (popularity / 10) : (client_queue' = client_queue + 1) & (ttl' = 1) & (responses_queue' = responses_queue + 1)
        +
        // Branch 2: Cache Miss
        (1 - (popularity / 10)) : (client_queue' = client_queue + 1) & (ttl' = 0) & (root_queue' = root_queue + 1);

    // 2. Query Root
    [DNS_Request_Resolver_To_Root] (client_queue > 0) & (root_queue > 0)
        -> (query_to_root_server' = true);

    // 3. Root Responds -> Move to Domain Queue
    [DNS_Response_Root_To_Resolver] (root_queue > 0) & (domain_queue < MAX_QUEUE)
        -> (root_queue' = root_queue - 1) & (domain_queue' = domain_queue + 1) & (query_to_root_server' = false);

    // 4. Query Domain (Opens Vulnerability Window)
    [DNS_Request_Resolver_To_Domain] domain_queue > 0 -> (query_to_domain_server' = true);

    // 5. Domain Responds (Closes Window)
    [DNS_Response_Domain_To_Resolver] (domain_queue > 0) & (responses_queue < MAX_QUEUE) & (correct_guess = false)
        -> (domain_queue' = domain_queue - 1) & (responses_queue' = responses_queue + 1) & (query_to_domain_server' = false)
        & (answer_from_domain_received' = true);

    // 6. Reply to Client
    [DNS_Response_Resolver_To_Client] (responses_queue > 0) & (client_queue > 0) -> (client_queue' = client_queue - 1)
        & (responses_queue' = responses_queue - 1);
    // -- Attack Race Condition
    [Guess] (correct_guess = false) & (query_to_domain_server = true) & (ctc_flag_active = false) ->
        1 / (query_id_range * port_id_range) : (correct_guess'=true)
        +
        ((query_id_range * port_id_range) - 1) / (query_id_range * port_id_range): (correct_guess'=false);

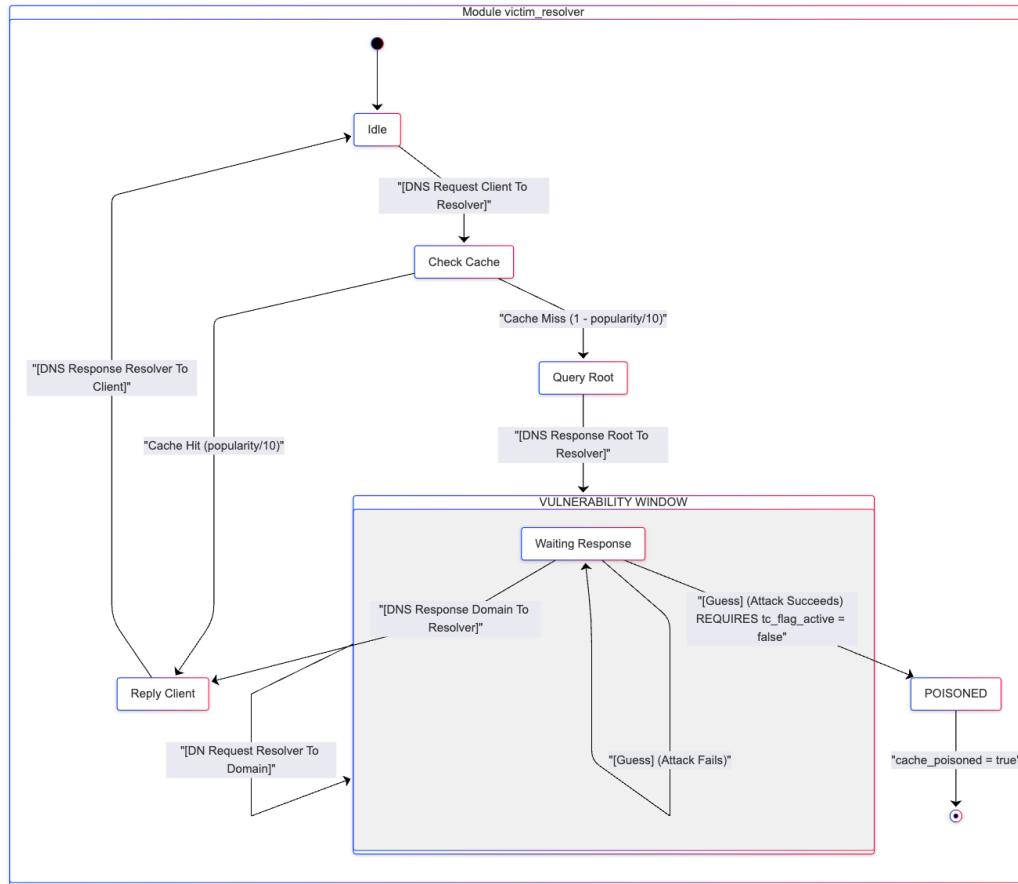
    // -- Sink State
    [] cache_poisoned = true -> true;
endmodule

```

Difference with previous  
Victim Resolver module



# PRISM Model of Kaminsky DNS Cache-Poisoning Attack with DNS-CPM Mitigation



1343

# Property verification and experiments

P = ? [F cache\_poisoned]

## Property Verification

**Property:**

P=? [ F cache\_poisoned ]

**Defined constants:**

NUMBER\_OF\_URL\_REQUESTS=1,popularity=0,port\_id\_range=1,guess=300,authoritative\_dns\_workload=150,benign\_noise\_rate=1000

**Method:**

Verification

**Result (probability):**

7.567437471068218E-5 (value in the initial state)



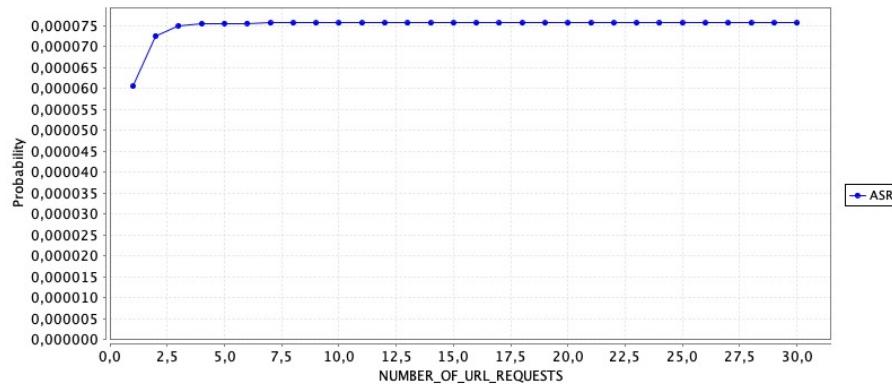
With only Query ID randomization (no Port randomization), a high guess rate, and a high workload on the Authoritative DNS Server, the attacker has a **0.0076%** of ***probability of poisoning the cache*** of the resolver



# Property verification and experiments

**P = ? [F cache\_poisoned]**

## Experiments



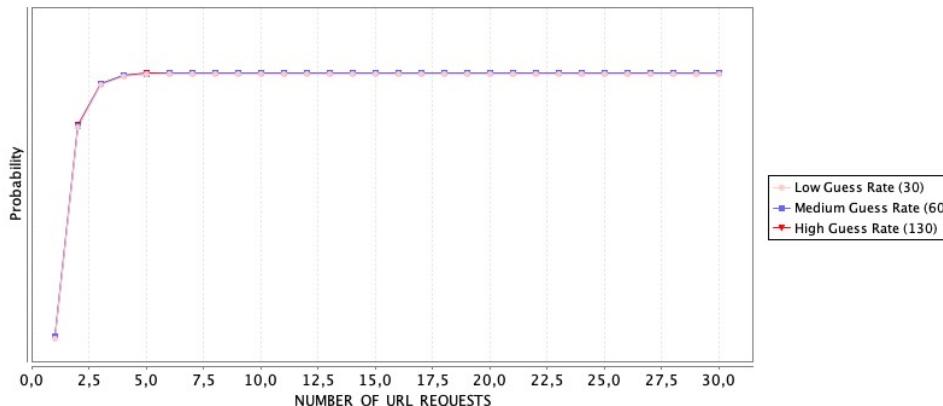
We plot the probability of attack success as a function of *number\_of\_url\_requests*, with *popularity*=2, *port\_id\_range*=1, *guess*=300, *authoritative\_dns\_workload*=35, *tau*=5, and *benign\_noise\_rate*=1000.

The probability of attack success changes for values of *number\_of\_url\_requests* varying from 1 to 5. After that, probability of success is almost constant.



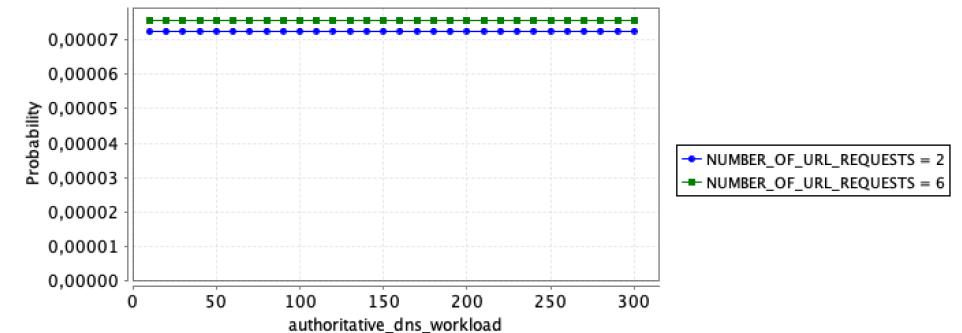
# Property verification and experiments

## Replication of the experiments done for model without DNS-CPM Mitigation



For three values of parameter *guess* (*low*=30, *medium*=60, *high*=130), we vary *number\_of\_url\_requests* from 1 to 30; *authoritative\_dns\_workload*=50, *port\_id\_range*=1, *tau*=5, and *benign\_noise\_rate*=1000.

The attack probability increases by increasing the number of URL-resolution requests until *number\_of\_url\_requests*=5. From that point on, attack probability remains constant.

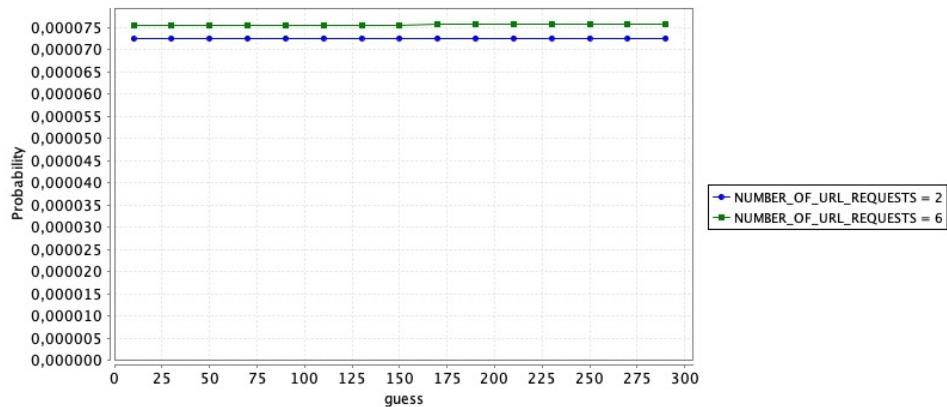


For two values of parameter *number\_of\_url\_requests* (2 and 6), we vary *authoritative\_dns\_workload* from 1 to 300; *guess*=50 and *port\_id\_range*=1, *tau*=5, and *benign\_noise\_rate*=1000.

Result is coherent with previous graph.



## Property verification and experiments

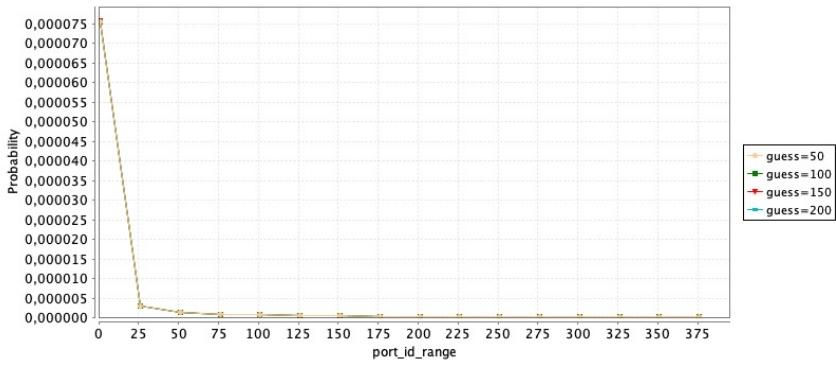
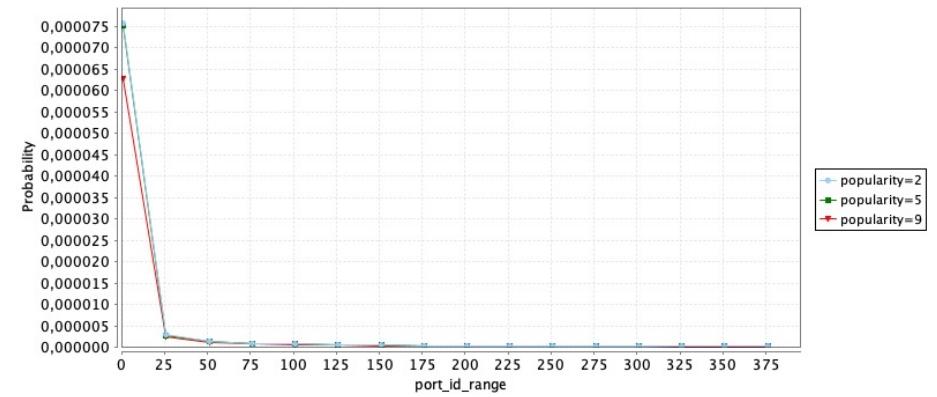
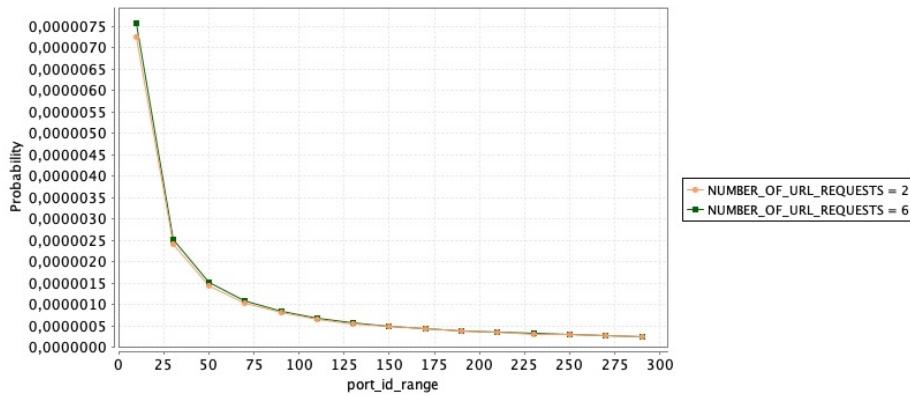


For two values of parameter *number\_of\_url\_requests* (2 and 6), we vary *guess* from 1 to 300; *authoritative\_dns\_workload* = 150, *port\_id\_range* = 1, *tau* = 5, *benign\_noise\_rate* = 1000.

**Observation:** we evince that the only parameter influencing attack success probability is the number of URL requests the attacker implements for the target domain.  
In particular, the attack probability changes for  $1 \leq \text{number\_of\_url\_requests} \leq 5$ .



# Property verification and experiments



In all the three experiments we can see that as soon as Port randomization is introduced, success probability of the attack goes to zero in a very fast way.



1343

# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack

## Recall DNS Cache-Poisoning attack exploiting fragmentation

- Attacker sends a series of 2nd fragment, with different IPIDs, containing the name-IP mapping to poison in resolver's cache → *larger attack window*
- Attacker sends a DNS Request to Resolver, which will force Authoritative DNS Server to respond with a fragmented IP packet
- No Port number to be guessed
- No Query ID to be guessed
- Only IPID to be guessed



# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack

## Modules

 **Intruder Machine**  
Intruder's PC that requests resolution of a URL

 **Victim Resolver**  
Local DNS Server that responds to the client's requests

**Root DNS Server**  
DNS Server that knows the IP address of the Authoritative DNS Server

**Domain Server**  
Authoritative DNS Server of the requested domain

 **Intruder Server**  
Intruder's DNS Server trying to guess correct Query ID (and Port number)

- Modules are the *same* of S type attack
- **Different behaviour** of *Intruder Server*, *Victim Resolver*, and *Domain Server*
- **Different condition** for attack success



# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack

```
// --- CONSTANTS ---
// Configuration
const int NUMBER_OF_URL_REQUESTS;
const int MAX_QUEUE = NUMBER_OF_URL_REQUESTS + 1;

// DNS Parameters
const double popularity;
const int ipid_range = 65536;

// Rates
const double guess;
const double requests_rate = 1.0;
const double authoritative_dns_workload;

// --- FORMULAS ---
formula cache_poisoned = (correct_guess = true & first_fragment_arrived = true) ; // Attack Status
```



Cache is **poisoned** (attack success condition) when attacker **correctly guesses 2nd fragment** and the **legitimate 1st fragment** from Authoritative DNS Server **arrives**



# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack

```
module intruder_machine
    trials : [0..NUMBER_OF_URL_REQUESTS] init 0;
    resolver_replies : [0..NUMBER_OF_URL_REQUESTS] init 0;

    // -- Transitions
    // 1. Send Request: Initiates resolution.
    [DNS_Request_Client_To_Resolver] trials < NUMBER_OF_URL_REQUESTS -> requests_rate : (trials' = trials + 1);

    // 2. Receive Reply: Marks end of query cycle.
    [DNS_Response_Resolver_To_Client] resolver_replies < trials -> (resolver_replies' = resolver_replies + 1);

    // 3. Termination
    [] resolver_replies = NUMBER_OF_URL_REQUESTS -> true;
endmodule
```

**Note:** all the requests made by the intruder to the resolver imply responses that need to be fragmented



# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack

```
module intruder_server
    is_state : [0..1] init 0;

    // -- Transitions
    // 1. Start the attack
    [DNS_Request_Client_To_Resolver] is_state = 0 -> (is_state' = 1);

    // 2. Firing Loop
    [Second_Fragment_Guess] is_state = 1 -> guess : (is_state' = 1);

    // 3. Stop the attack
    [DNS_Response_Domain_To_Resolver_Frag2] is_state = 1 -> (is_state' = 0);
endmodule
```

- Attacker **starts sending 2nd fragment guesses** together with the request being sent to the resolver → We model this behaviour by synchronizing the beginning of the attack with the Intruder Machine sending the request to the resolver.
- Attacker stops if the legitimate 2nd fragment arrives to the resolver.



# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack

```
module victim_resolver
    ttl : [0..2] init 2; // 0=Miss, 1=Hit, 2=Init
    client_queue : [0..MAX_QUEUE] init 0;
    root_queue : [0..MAX_QUEUE] init 0;
    domain_queue : [0..MAX_QUEUE] init 0;
    responses_queue : [0..MAX_QUEUE] init 0;
    first_fragment_arrived : bool init false;
    second_fragment_arrived : bool init false;

    // -- Variables
    query_to_root_server: bool init false;
    query_to_domain_server: bool init false;
    answer_from_domain_received: bool init false;
    correct_guess: bool init false;

    // -- Transitions
    // 1. Client Request (Merged Hit/Miss Logic)
    [DNS_Request_Client_To_Resolver] (client_queue < MAX_QUEUE) & (responses_queue < MAX_QUEUE) & (root_queue < MAX_QUEUE) ->
        // Branch 1: Cache Hit
        (popularity / 10) : (client_queue' = client_queue + 1) & (ttl' = 1) & (responses_queue' = responses_queue + 1)
        +
        // Branch 2: Cache Miss
        (1 - (popularity / 10)) : (client_queue' = client_queue + 1) & (ttl' = 0) & (root_queue' = root_queue + 1);

    // 2. Query Root
    [DNS_Request_Resolver_To_Root] (client_queue > 0) & (root_queue > 0)
        -> (query_to_root_server' = true);

    // 3. Root Responds -> Move to Domain Queue
    [DNS_Response_Root_To_Resolver] (root_queue > 0) & (domain_queue < MAX_QUEUE)
        -> (root_queue' = root_queue - 1) & (domain_queue' = domain_queue + 1) & (query_to_root_server' = false);

    // 4. Query Domain (Opens Vulnerability Window)
    [DNS_Request_Resolver_To_Domain] domain_queue > 0 -> (query_to_domain_server' = true);

    // 5. Domain Responds with first fragment
    [DNS_Response_Domain_To_Resolver_Frag1] (domain_queue > 0) -> (first_fragment_arrived' = true);
```



# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack

```
// 6. Domain Responds with second fragment (Closes Window)
[DNS_Response_Domain_To_Resolver_Frag2] (responses_queue < MAX_QUEUE) & (domain_queue > 0) & (correct_guess = false)
-> (domain_queue' = domain_queue - 1) & (responses_queue' = responses_queue + 1) & (query_to_domain_server' = false)
& (answer_from_domain_received' = true) & (second_fragment_arrived' = true);

// 7. Reply to Client
[DNS_Response_Resolver_To_Client] (responses_queue > 0) & (client_queue > 0) -> (client_queue' = client_queue - 1) & (responses_queue' = responses_queue - 1)
& (first_fragment_arrived' = false) & (second_fragment_arrived' = false);

// -- Attack Race Condition
[Second_Fragment_Guess] (correct_guess = false) & (second_fragment_arrived = false) ->
1 / (ipid_range) : (correct_guess'=true)
+
((ipid_range - 1) / (ipid_range)) : (correct_guess'=false);

// -- Sink State
[] cache_poisoned = true -> true;
endmodule
```

- The attack window closes when resolver receives the legitimate 2nd fragment from Authoritative DNS Server
- The rate at which the attacker correctly guesses the 2nd fragment is given by:

$$guess * \frac{1}{ipid\_range} = guess * \frac{1}{2^{16}}$$



# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack

```
module root_server
    root_state : [0..1] init 0;

    // -- Transitions
    // 1. Receive Request: A query arrives from the Victim Resolver.
    [DNS_Request_Resolver_To_Root] root_state = 0 -> (root_state' = 1);

    // 2. Send Reply: The Root Server immediately responds, moving the resolver to the next step (Domain Server).
    [DNS_Response_Root_To_Resolver] root_state = 1 -> (root_state' = 0);
endmodule
```



# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack

```
module target_domain_server
    // -- States
    domain_state : [0..2] init 0;

    // -- Transitions
    // 1. Receive Request: Opens Race Window
    [DNS_Request_Resolver_To_Domain] domain_state = 0 -> (domain_state' = 1);

    // 2. Legitimate Response: Closes Race Window
    [DNS_Response_Domain_To_Resolver_Frag1] domain_state = 1 -> 1 / (authoritative_dns_workload) : (domain_state' = 2);
    [DNS_Response_Domain_To_Resolver_Frag2] domain_state = 2 -> 1 / (authoritative_dns_workload) : (domain_state' = 0);

endmodule
```

Since DNS response needs to be fragmented, Authoritative DNS Server sends:

1. 1st fragment
2. 2nd fragment

Of course, these operations are done together with management of *other legitimate requests arriving to the Authoritative DNS Server*. This gives the attacker higher opportunity to guess the correct 2nd fragment.



# Property verification and experiments

P = ? [F cache\_poisoned]

## Property Verification

**Property:**

P=? [ F cache\_poisoned ]

**Defined constants:**

NUMBER\_OF\_URL\_REQUESTS=1,popularity=0,guess=300,authoritative\_dns\_workload=150

**Method:**

Verification

**Result (probability):**

0.6532616691133615 (value in the initial state)



With  $S_{frag}$  attack, a high guess rate, and a high workload on the Authoritative DNS Server, the attacker has a **65%** of **probability of poisoning the cache** of the resolver



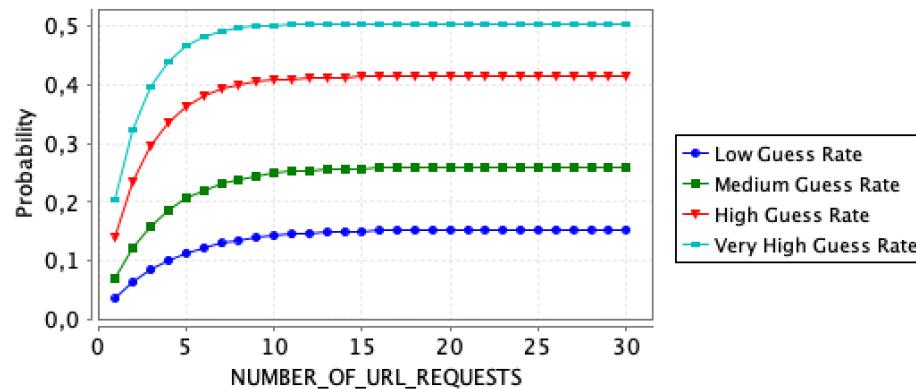
This attack has a higher success probability than Kaminsky Attack



# Property verification and experiments

P = ? [F cache\_poisoned]

## Experiments



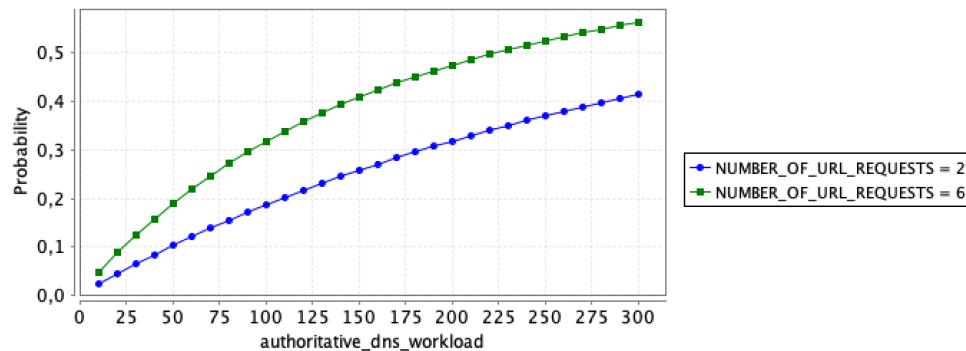
For different values of *guess rate* (low=30, medium=60, high=130, very high=200), *authoritative\_dns\_workload*=50, and *popularity*=2, we varied *number\_of\_url\_requests* from 1 to 30.

Result shows that probability of attack success increases for *number\_of\_url\_requests* between 1 and 15, and for increasing values of *guess rate*.



# Property verification and experiments

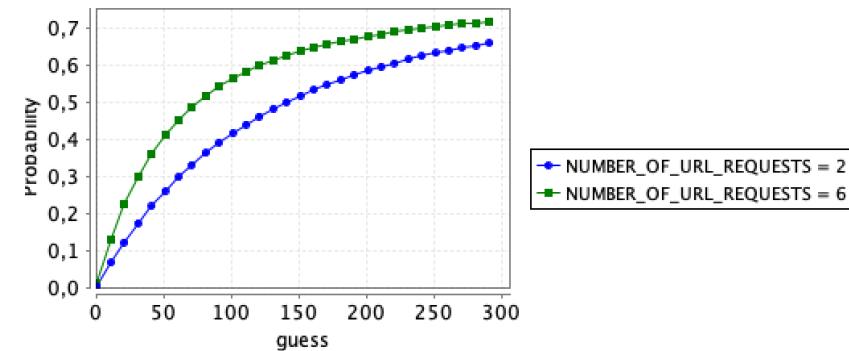
P = ? [F cache\_poisoned]



For two values of parameter *number\_of\_url\_requests* (2 and 6), we vary *authoritative\_dns\_workload* from 10 to 300; *guess*=50, and *popularity*=2.

As expected, the higher the workload on Authoritative DNS Server, the higher the probability of attack success.

Moreover, as we've already seen in previous experiment, higher values of *number\_of\_url\_requests* lead to higher probabilities.



For two values of parameter *number\_of\_url\_requests* (2 and 6), we vary *guess* from 1 to 300; *authoritative\_dns\_workload* =150, and *popularity*=2. As expected, increasing rate *guess* increases the probability of attack success.

# DNS-CPM Mitigation (Rℓ2)



## Detection Module (Rℓ2)



Fragments are identified as potential indicators for a DNS Cache-Poisoning Attack.

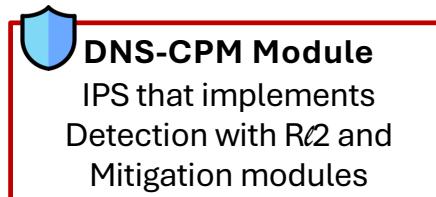
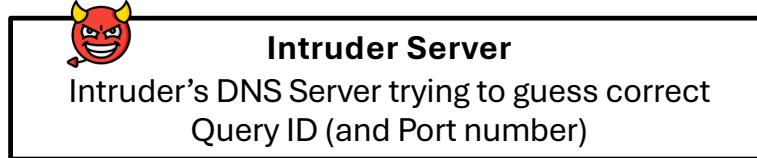
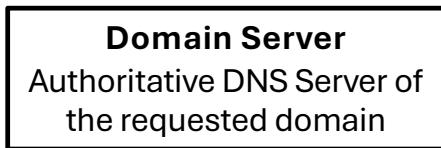
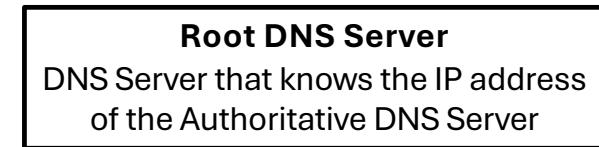
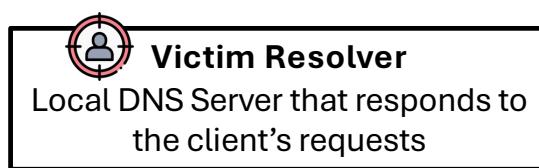
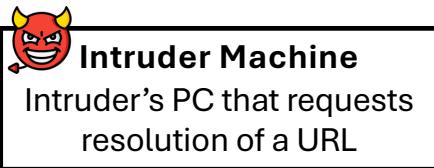


When the 1st fragment of a response is detected, the response is flagged as a **potential poisoning attack** and passed to the Mitigation Module



# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack with DNS-CPM Mitigation

## Modules



# PRISM Model of S<sub>Frag</sub> DNS Cache-Poisoning Attack with DNS-CPM Mitigation

```
module dns_cpm_module

    // Signal to the Mitigation Module (True if threshold reached)
    detection_signal : bool init false;

    // 1. Detection of a Possible Attack
    [DNS_Response_Domain_To_Resolver_Frag1] true -> (detection_signal' = true); → Detection module kicks in mitigation when a 1st fragment is detected

endmodule

formula tc_flag_active = (detection_signal = true); // Mitigation Defense
formula cache_poisoned = (correct_guess = true & first_fragment_arrived = true & tc_flag_active = false) ; // Attack Status
```

When mitigation starts, cache cannot be poisoned even if correct 2nd fragment and legitimate 1st fragment has arrived to the resolver



# Property verification

P = ? [F cache\_poisoned]

## Property Verification

**Property:**  
P=? [ F cache\_poisoned ]

**Defined constants:**  
NUMBER\_OF\_URL\_REQUESTS=1,popularity=0,guess=300,authoritative\_dns\_workload=150

**Method:**  
Verification

**Result (probability):**  
0.0 (value in the initial state)



With DNS-CPM R<sub>2</sub>, 0% of **probability of poisoning the cache** of the resolver



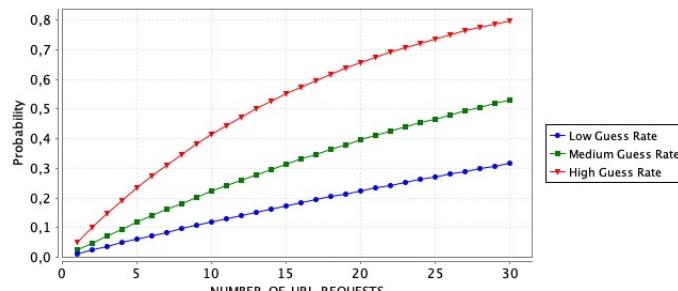
This is the result we expected. To poison the cache, both legitimate 1st fragment and correct forged 2nd fragment must arrive to resolver. However, **as soon as the 1st legitimate fragment arrives, detection module activates mitigation** which moves conversation to DNS over TCP



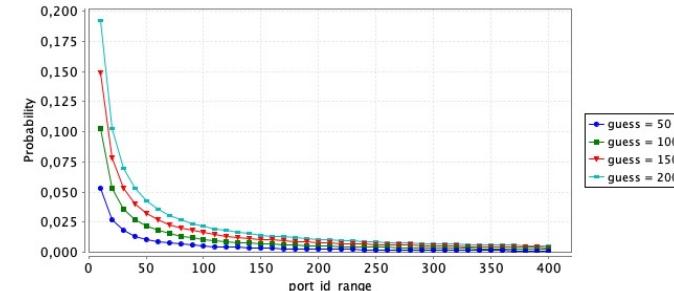
# Summarizing

## Kaminsky DNS Cache-Poisoning Attack:

- Probability of attack success is strictly related to *guess rate*, *number of URL requests* sent by the attacker to the resolver, and the *workload on Authoritative DNS Server*.
- Attack become more difficult when both Query ID randomization and Port Number randomization are employed.



Randomization of Query ID only



Randomization of Query ID and Port number

## DNS-CPM R1:

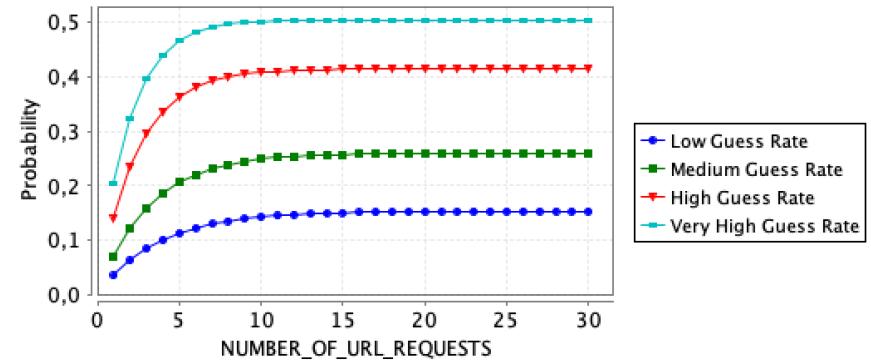
- Strongly reduces attack success probability to 0.0076% with only Query ID randomization.
- Notice that once detection triggers the mitigation module, session switches to **DNS-over-TCP**



# Summarizing

## DNS Cache-Poisoning Attack based on Fragmentation:

- With this type of attack, the intruder can overcome the need of guessing Query ID and Port number.
- It has only to guess a 16-bit IPID.
- To do that, it has a larger window since it starts sending forged 2nd fragment immediately.
- As we've seen, probability of attack is higher than 65%



## DNS-CPM R<sub>2</sub>:

- Detection with R<sub>2</sub> brings probability of attack success to 0%



# Observations and conclusions

## Out-of-Bailwick DNS Cache-Poisoning Attack

- Variant of Kaminsky attack
- Approach is the same, but attacker inserts information in Additional Records trying to poison more domain names.
- For this reason, we didn't model it.

## DNS-CPM R@3:

- It verifies that each record in the DNS response belongs to the same domain (or subdomain) requested in the original query.
- If a record does not belong to the same domain, the packet is immediately classified as suspicious and passed to the mitigation module.
- At that point, session will pass to DNS over TCP, and the attacker cannot poison the resolver's cache.

### DNS-CPM Mitigation Limit

Not compatible with the small percentage of resolvers (approx. 2.67%) that ignore the TC flag and do not retry over TCP.



# References

---

[Formal Analysis of the Kaminsky DNS Cache-Poisoning Attack Using Probabilistic Model Checking](#)

---

[DNS-CPM: DNS Cache Poisoning Mitigation](#)

---

Network Security course notes

---

Computational Models for Complex Systems course notes

---

[PRISM Model Checker](#)

---

[Diagrams Tool: Mermaid JS](#)

---

